

# Lecture 17: Authentication and Access Control

Stephen Huang

## Content

1. Authentication
2. Access Control (Authorization)
3. Unix Access Control

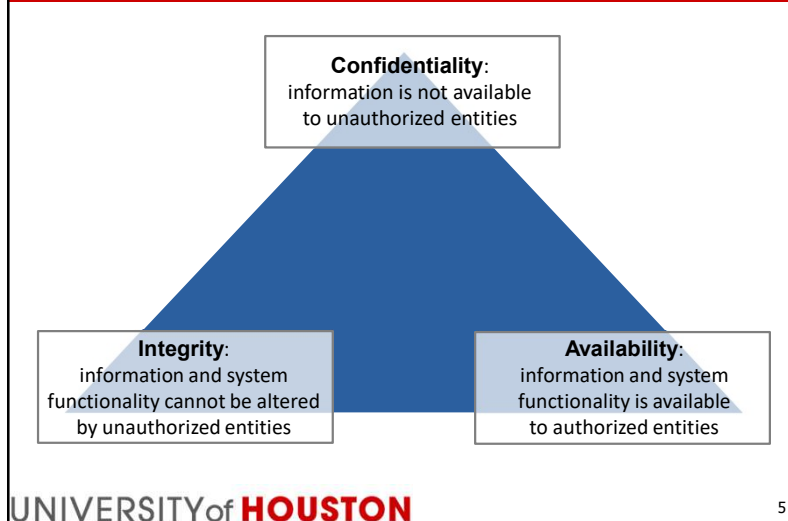
## Vulnerabilities



## Vulnerabilities



## Reminder: Security Objectives



5

## Vulnerabilities



6

## 1. Authentication

**Authentication:** reliably verifying the identity of someone or something

- computer authenticates another computer
- computer authenticates a user
- Typical methods of computer authentication
  - cryptography-based  
(*example:* using the Kerberos protocol or using public-key certificates)
  - address-based  
(*example:* identifying a computer based on its IP address)
- Types of authentication
  - one-way authentication or mutual authentication
  - one-time or establishing a session (*e.g.*, combined with key exchange)

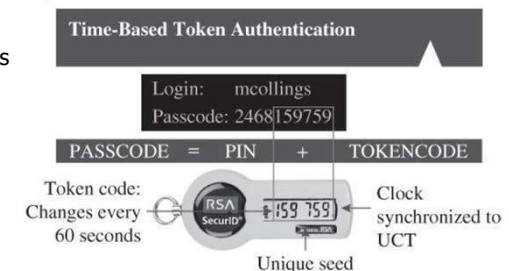
UNIVERSITY of HOUSTON

7

## Authentication Mechanisms

Authentication is based on something you know, are, or have.

- Something you know: password, PIN, phrase, facts
- Something you are: Biometrics such as fingerprint, voice, faces
- Something you have:
  - Static tokens such as ID badges, physical keys, chips (on credit card),
  - Dynamic tokens such as SecureID Token



UNIVERSITY of HOUSTON

8

7

8

## User Authentication

- Types of user authentication factors
  - knowledge: some secret known only by the user (e.g., password)
  - ownership: some physical object possessed by the user (e.g., bank card)
  - inherence: some physical characteristic of the user (e.g., fingerprint)
- Password-based user authentication
  - typical form of knowledge-based authentication
  - verifier stores the password in a database or file
  - often combined with cryptography-based approaches to protect the password from eavesdropping
  - password must be easy to remember but hard to guess

## Password-Based Authentication

- *Problem*: easy-to-remember passwords are weak
  - Miller's law: the number of objects an average human can hold in working memory is  $7 \pm 2$ .
  - Published in 1956 by cognitive psychologist George A. Miller ("The Magical Number Seven, Plus or Minus Two")
  - The length of passwords that users can easily remember (i.e., not write down somewhere) is very limited
    - *example*: 8 alphanumeric characters  $\rightarrow 36^8$  possibilities  $\sim 2^{48}$  possibilities
      - $\rightarrow$  brute-force guessing may be computationally feasible
    - most popular passwords of 2019 (according to *SplashData*):
      1. 123456
      2. 123456789
      3. qwerty
      4. password
      5. 1234567

## Brute-force Attack

- Brute-force attack: password guessing
  - online: attacker must rely on the verifier to test the correctness of a password
    - $\rightarrow$  verifier can limit the number of attempts (e.g., number of unsuccessful login)
  - offline: attacker can test the correctness of a password on its own

## Password Storage

- Cleartext passwords are insecure
  - system administrators (and other local users) may easily read passwords
  - attackers who have compromised a system may be able to read passwords
- Example incident: Yahoo data breach
  - in September 2016, Yahoo announced that hackers breached its system sometime in late 2014
  - hackers accessed personal information (e.g., names, e-mail addresses, dates of birth, ...) associated with 500 million Yahoo! user accounts
- Users tend to reuse passwords
  - $\rightarrow$  breach may affect other systems as well

## Storing Hashed Passwords

- Store the cryptographic hash of the password
  - during authentication, the user enters the plaintext password, and the verifier computes its hash and compares it with the stored hash
  - The attacker can perform offline guessing to recover the plaintext password
- Example: Unix systems
  - on modern systems, hashed passwords are stored in `/etc/shadow`, which can be read only by the root user
  - non-sensitive information is stored in the file `/etc/passwd`, which is readable by all local users
- Brute-forcing multiple hashed passwords
  - first, precompute a table of [password, hash] values for possible passwords
  - second, for each hashed password, look up the precomputed hash value

## Precomputed Hash Chains

- Lookup Table of all possible passwords.

```
4a7d1ed414474e4033ac29ccb8653d9b:0000
25bbdcd06c32d477f7fa1c3e4a91b032:0001
fcd04e26e900e94b9ed6dd604fed2b64:0002
...
fa246d0262c3925617b0c72bb20eeb1d:9999
```

- Attacker's problem: The list of possible passwords is too long, which means prohibitive space requirement for storing precomputed hashes.
- Precomputed hash chain: trading off space for running time.

## Building Hash Chains

- Reduction function  $R$ : maps a hash value to a possible password (**not** the actual inverse of the hash function)
- Choose a random set of initial passwords (e.g., aaaaaa, bbbbbb, ...)
- For each password,
  - compute a chain of passwords and hash values by alternating between using the hash and reduction functions, e.g.:

aaaaaa  $\xrightarrow{H}$  281DAF40  $\xrightarrow{R}$  sgfnvd  $\xrightarrow{H}$  920ECF10  $\xrightarrow{R}$  kiebgt

- store the initial and final passwords (e.g., (aaaaaa, kiebgt)) in a table.

- The reduce function is essentially a deterministic pseudo-random password generator.


## Recovering a Password

- To recover a hashed password, start building a chain from the hash value, and test if any of the resulting passwords is among the stored final passwords
  - e.g., if hash is 920ECF10, then build chain to kiebgt
- When a match is found, the correct password can be recovered from the chain
  - e.g., from stored initial password aaaaaa, we can build chain to sgfnvd

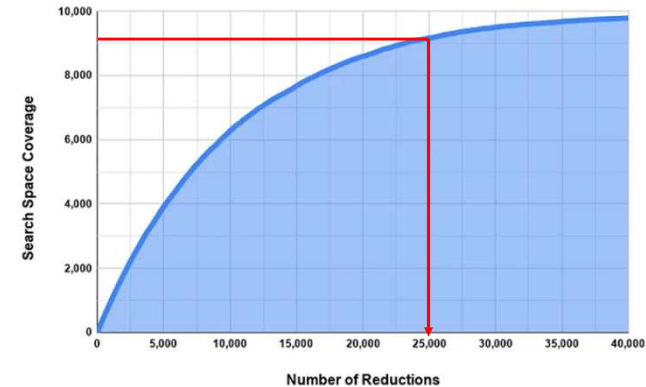
aaaaaa  $\xrightarrow{H}$  281DAF40  $\xrightarrow{R}$  sgfnvd  $\xrightarrow{H}$  920ECF10  $\xrightarrow{R}$  kiebgt

## Rainbow Table

- Hash chain limitation: when hash or password values collide, the remainder of the chains (including the final value) are the same
  - number of usable chains is limited
- Rainbow table
  - to prevent merging chains, use a sequence of reduction functions  $R_1, \dots, R_k$



## Probability of a collision



## Rainbow Table

- When one value collides, all subsequent values collide, too. Merging.
- It gets harder and harder to avoid a collision.
- We cannot efficiently detect the collision at generation time since that would require holding on to the full contents of the previous chains.
- The solution to merging is to use a sequence of reduce functions  $R_i$  in Step  $i$ .
- A practical implementation of a sequence of reduce function is to add a second parameter of "salt".
- Merges are still possible, but only if the collision "lines up".

## Salting

- Before hashing a password, mix it with a salt value
  - both when the password is set and during verification
  - verifier stores: username, salt,  $H(\text{password} + \text{salt})$
- Salt value
  - randomly generated for each user account
  - may be stored in plaintext by the verifier
- Salt values do not have to be memorized → strong randomness
  - prevents precomputing hashes since the attacker cannot consider all possible salt values (different salt values require different precomputation)
  - also hides identical passwords, which would result in identical hashes
- However, it does not make guessing a single password harder (assuming that the attacker knows the salt)



## Salting Issues

- Example: 2012 LinkedIn hack
  - on June 5, LinkedIn was breached, and around 6.5 million hashed passwords were stolen
  - on June 6, a large number of recovered plaintext passwords were posted online
  - in May 2016, it was discovered that an additional 100 million might have been compromised in the incident
  - Major weaknesses:
    - passwords were not salted before hashing
    - passwords were hashed using SHA-1
- SHA-1 is relatively easy to compute
  - brute-force guessing is relatively fast



## Rainbow Tables

- Rainbow tables aren't popular anymore.
- Brute-force attacks aren't that effective.
- Secured hashing is immune since it uses salting.
- The best rainbow tables publicly available only go up to 8 characters for a full character set.
- The average password length is over 9 characters,
- Other methods are better. Wordlist attacks with manipulation rules are far more effective at getting actual user-picked passwords.

## Multi-Factor Authentication


- User is authenticated only after passing multiple independent authentication mechanisms
  - typically, each mechanism is built on a different type of factor (e.g., knowledge + possession), so it is independent of the other mechanisms
  - The attacker must circumvent all authentication mechanisms to succeed
- Possession factors
  - disconnected token: not connected to the client computer, typically, the user manually enters authentication data displayed by the token
  - connected token: physically connected to the client computer (e.g., USB token)
- Inherence factors
  - includes fingerprint, face, voice, or iris recognition



example:  
RSA SecurID token

## 2. Access Control - Authorization

- Access control (i.e., authorization): approving or rejecting access requests.
- Abstractions
  - subjects: entities that can perform actions on the system
  - objects: resources to which access must be controlled
- Control access to objects based on a policy



## Discretionary Access Control (DAC)

- Allows access rights to be propagated at the subjects' discretion
- Often implemented using the notion of owner
  - every object has an owner subject, who can set the permissions for that object
- Used by popular operating systems (e.g., Unix and Windows)


**Problem:** non-malicious users are not necessarily trustworthy

- phishing: subjects may be tricked into propagating their access rights to malicious entities
  - malware: malicious code running with a subject's credentials can disclose or modify sensitive information
- large organizations working with sensitive data may need centralized control

## Mandatory Access Control (MAC)

- Restricts the access of subjects to objects based on a system-wide set of rules
  - system-wide rules are set by a central authority (e.g., system administrator)
  - policy is mandatory → users do not have full control over access to the resources that they create
- Traditionally used for implementing multilevel security
  - objects have security classifications (e.g., "Top Secret", "Secret")
  - subjects have security clearances
- Available in some form on many modern operating systems
  - SELinux and AppArmor for Linux, and Mandatory Integrity Control for Windows
- May be combined with DAC: grant access only if both DAC and MAC permit the access


## Access Control Models



- Access control list (ACL): list permissions for each object
  - for each object, list pairs of [subject, access right]
- Role-based access control (RBAC): row oriented
  - create a set of roles (e.g., based on real-world job functions), and assign a role (or roles) to each subject
  - for each role, list pairs of [object, access right]

## Bell-LaPadula Model

- Developed by D. Bell and L. LaPadula in the 1970s for enforcing access control in government and military applications
- Multilevel security (e.g., "Top Secret", "Secret")
  - objects have security classifications
  - subjects have security clearances
- Focuses on confidentiality
- Rules
  1. **simple security property:** subjects cannot read objects at a higher security level
  2. **\*-property:** subjects cannot write to objects at a lower security level
  3. **discretionary security property:** use an access matrix to specify DAC
- Information may be transferred from a higher level to a lower level by trusted subjects





### 3. Unix Access Control

- Basic Concepts
  - user: has a unique UID (special UID = 0 for root user)
  - group (collection of multiple users): has a unique GID
- Access control abstraction
  - subject = process
    - has an effective UID and GID (as well as real and saved UIDs and GIDs)
  - object = file
    - has an owner (UID) and a group (GID), typically inherited from the process that created the file
    - almost everything is a file on a Unix system (regular files, directories, devices, Unix domain sockets, ...)

### Example

```

student@server:/stuff/$ ls -l
-rw--xrw- 1 TA      web_sec  4096 Apr 21 finalExam.txt
d-wx-w-r-- 3 instructor web_sec 4096 Apr 17 grading
d---rwsrwt 7 instructor teachers 4096 Apr 16 security_memes
  
```

User      Group      Object

### Unix Access Control: Permission

- Each file has 12 permission bits
  - **read**, **write**, and **execute** permission for owner, group, and others
  - set user ID (setuid), set group ID (setgid), sticky bits
- When a process wants to read/write/execute a file,
  1. if effective UID = file owner → use read/write/execute permission for owner
  2. else if effective GID = file group → use read/write/execute permission for group
  3. else → use read/write/execute permission for others
- For directories,
  - read means listing the contents of the directory
  - write means creating, renaming, and deleting files in the directory
  - execute means accessing the files (and directories) within the directory (must also have execute permission on all the parent directories)

### Example

```

student@server:/stuff/$ ls -l
-rw--xrw- 1 TA      web_sec  4096 Apr 21 finalExam.txt
d-wx-w-r-- 3 instructor web_sec 4096 Apr 17 grading
d---rwsrwt 7 instructor teachers 4096 Apr 16 security_memes
  
```

User rwx      Group rwx      Other rwt      User      Group      Object

The sticky bit should have been called the "restricted deletion (by owner) bit".



## Sticky, Set UID, and Set GID Bits

- Sticky bit
  - when set on a directory, files within that directory can be renamed or deleted only by their owners, the directory owner, or a superuser
  - for example, sticky bit is typically used on the `/tmp` directory
- Set UID bit
  - when set on an executable file, the effective UID of a process executing the file is set to the file owner UID
  - for example, set UID bit is typically used on the `passwd` command
- Set GID bit
  - when set on an executable file, the effective GID of a process executing the file is set to the file group GID
  - when set on a directory, new files created within will inherit the GID of the directory

## Example

```

student@server:/stuff/$ ls -l
-rw-r--rwx 1 TA      web_sec  4096 Apr 21 finalExam.txt
d-wx-wx-- 3 instructor web_sec 4096 Apr 17 grading
d---rwsrwt 7 instructor teachers 4096 Apr 16 security_memes
  
```

Annotations in the image:

- User:** Points to the first character of the permission string (e.g., 'r' in 'rw-r--rwx').
- Group:** Points to the third character of the permission string (e.g., '-' in 'rw-r--rwx').
- Object:** Points to the file name (e.g., 'finalExam.txt').

## Unix Access Control Conclusion

- Processes running with `setuid/setgid`
  - effective UID/GID is the UID/GID of the executable file, while the real UID/GID is the UID/GID of the parent process
- Changing the owner or group of a file
  - only a superuser can change ownership
  - only a member of a group can change the group of a file to that group
- Traditional Unix access control is DAC / ACL
- Some Unix versions offer support for other policies
  - SELinux (Security Enhanced Linux): support for RBAC or MAC
  - Oracle Solaris: support for RBAC

## Example

```

cooluser@LAPTOP-5V55HON5:/topsecretfolder$ ls -l
total 8
drwxrwx--x 2 root root 4096 Oct 18 19:34 secret1
drwxrwx-- 2 root root 4096 Oct 18 19:31 secret2
cooluser@LAPTOP-5V55HON5:/topsecretfolder$

cooluser@LAPTOP-5V55HON5:/topsecretfolder$ ls secret1
pepperNeggRecipe.txt
cooluser@LAPTOP-5V55HON5:/topsecretfolder2$

cooluser@LAPTOP-5V55HON5:/topsecretfolder$ ls secret2
ls: cannot open directory 'secret2': Permission denied
cooluser@LAPTOP-5V55HON5:/topsecretfolder3$

root@LAPTOP-5V55HON5:/topsecretfolder# ls secret2
GardinieraRecipe.txt
root@LAPTOP-5V55HON5:/topsecretfolder#
  
```

Annotations in the image:

- Root:** Points to the 'root' user/group in the first two lines of the first terminal output.

## Next Topic

- Authentication and Access Control
- Software Security