

Lecture 19: Other Vulnerabilities

Stephen Huang

Content

1. Integer Overflow
2. Input Validation Weaknesses
3. Format String Vulnerabilities
4. Race Conditions
5. Buffer Overflow Countermeasures
 1. Com-Time Prevention
 2. Run-Time Countermeasures

1. Integer Overflow

- An integer overflow occurs when you attempt to store inside an integer variable a value larger than the maximum value the variable can hold.
- The C standard defines this situation as undefined behavior (meaning that anything might happen).
- In practice, this usually translates to a wrap of the value if an unsigned integer was used and a change of the sign and value if a signed integer was used.
- Even some managed languages, such as Java and C#, are susceptible to integer overflow errors.

Integer Overflow

- Every integer type has a maximum value.
 - Maximum for 32-bit int is 2,147,483,647
 - The range for 8-bit byte is 0 to 255
- Example


```
unsigned char x = 100;
unsigned char y = 200;
unsigned char z = x + y;

# z == 44 instead of z == 300
```
- Affects most languages
 - even some managed languages, such as Java and C#, are susceptible to integer overflow errors

Integer Overflow Vulnerabilities

- Exploiting integer overflow errors
 - calculating indexes into arrays
 - calculating the amount of space to allocate for a buffer
 - checking whether an overflow could occur
- Example: // input was provided by the user

```
const unsigned short SIZE = 10000;
char buffer[SIZE];
unsigned short length = strlen(input);
if (length < SIZE) strcpy(buffer, input);

bool checkOverflow(unsigned short x, unsigned short y)
{
    if (x + y < x) return true;
    return false;
}
```

Does it work?

Incorrect check due to integer promotion

Integer Overflow Vulnerabilities

```
#include <stdio.h>
int main() {
    char a = 30, b = 40, c = 10;
    char d = (a * b) / c;
    printf ("%d ", d);
    return 0;
}
```

Output: 120

The range for a signed char For most C compilers is -128 to 127

Integer promotion occurs before the division.

2. Input Validation Weaknesses

- Sources of input
 - User-supplied files and terminal input
 - Command line arguments
 - Environment variables
 - Function calls from other modules
 - Network packets (web applications in detail later)
- Never trust your input
 - Specially crafted input can cause buffer overflows, integer overflows, ...
 - Input should always be validated
 - Lack of input validation can lead to software vulnerabilities



Command Injection

- Vulnerable code:

```
void sendEmailToUser(char *username) {
    char buffer[1024];
    sprintf(buffer,
        "mail -s 'Please do not hack us' %s@example.com", username);
    system(buffer);
}
```

- username = "user"

```
system("mail -s 'Please do not hack us' user@example.com");
```

- username = "foo@bar.com ; rm very_important_file ;"

```
system("mail -s 'Please do not hack us' foo@bar.com ;
    rm very_important_file ; @example.com")
```

could be any system command, which will be executed by the exploited process

3. Format String Vulnerabilities

- Uncontrolled format strings were discovered to be a vulnerability around 1989, first successful exploits were published around 2000.
- Format string
 - `printf("Security is very %s.", "important");`
 - `printf("The %s is %d!", "answer", 42);`
- Vulnerable functions
 - `sprintf`: writes to buffer
 - `fprintf`: writes to file
 - other members of the `printf` family (e.g., `snprintf`)
 - `printf`: used in the Linux kernel
 - other functions that use format strings (e.g., `syslog`)
- Format placeholders
 - `%s` – string (reference)
 - `%d` – number (output in decimal format)
 - `%x` – number (output in hexadecimal format)
 - ...

Safe printf()

```
#include <stdio.h>
void main(int argc, char **argv)
{
    // This line is safe
    printf("%s\n", argv[1]);
}
```

`./example "Hello World %s%s%s%s%s"`

- The `printf` will not interpret the `"%s%s%s%s%s"` in the input string, and the output will be: `"Hello World %s%s%s%s%s"`

Vulnerable printf()

```
#include <stdio.h>
void main(int argc, char **argv)
{
    // This line is vulnerable
    printf(argv[1]);
}

./example "Hello World %s%s%s%s%s"
```

- The `printf` will interpret the `%s%s%s%s%s` in the input string as a reference to string pointers, so it will try to interpret every `%s` as a pointer to a string, starting from the location of the buffer (probably on the Stack).
- At some point, it will get to an invalid address, and attempting to access it will cause the program to crash.

Format String Vulnerabilities

```
int main() {
    char string[1024];
    gets(string);
    printf("You wrote: %s\n",
          string);
}
```

```
$ ./program
test
You wrote: test
```

```
$ ./program
%x %x
You wrote: %x %x
```

```
int main() {
    char string[1024];
    gets(string);
    printf("You wrote: ");
    printf(string);
    printf("\n");
}
```

```
$ ./program
test
You wrote: test

$ ./program
%x %x
You wrote: fb7d6460a 7a3b64d06
```

`printf` tries to read parameter values for `%x` from the stack

Exploiting Format Strings

```
int main() {
    char string[1024];
    int secret = 42;

    gets(string);
    printf("You wrote: ");
    printf(string);
    printf("\n");
}

$ ./program
test
You wrote: test

$ ./program
%d
You wrote: 42
```

Actual stack setup for
printf

Stack setup assumed by
printf for %d

local variables of main	string	local variables of main	
	secret	parameters of printf	integer
parameter of printf	pointer to string		pointer to string
return address to main			
saved frame pointer			
local variables of printf			

Reading Memory

```
int main() {
    char *secret = malloc(1024);
    char string[1024];

    load_secret_info(secret);
    gets(string);
    printf(string);
}
```

Attack string:

%d%d...%d %s
 256 x 4 bytes Points to secret info

Stack setup assumed by printf
for %d...%d%s

local variables of main	pointer secret buffer string	parameters of printf	data for %s data for %d...%d
parameter of printf	pointer to string		pointer to string
return address to main			
saved frame pointer			
local variables of printf			

alternatively, attacker
can supply a memory
address for %s within
the string
→ attacker can trick
the process into
printing the contents
of memory at the
chosen location

Writing to Memory Using Format

- Special placeholder: %n
 - argument must be a pointer to a signed integer, where the number of characters printed so far will be written
 - example: `printf("foobar%n", &x)` writes the value 6 to variable x
- Providing a pointer
 - pointer can be part of the malicious string, since it is stored on the stack
 - attacker can use %x or %d placeholders to reach it
- Controlling the value
 - attacker can control the length of the string
- Arbitrary code execution
 - embed shellcode in the string
 - overwrite return address so that it points to the shellcode

Arbitrary Code Execution

```
int func() {
    char string[1024];
    gets(string);
    printf(string);
}
```

Points to the location of
the return address on
the stack

- Attack string:

%99d%99d...%99d %n <<address>><<shellcode>>

Two purposes:

- going upwards in the stack until we reach <<address>>
- controlling the length of the string → value written to the return address

attacker can overwrite
the return address with
a value that points to
the <<shellcode>>

Arbitrary Code Execution

local variable of func (i.e., string)	%99d%99d...%99d %n <<address>> <<shellcode>>	local variables of func	%99d%99d...%99d %n
parameter of printf	pointer to string	parameters of printf	data for %n data for %99d...%99d pointer to string
return address to func			
saved frame pointer			
local variables of printf			

Format String Vulnerabilities

- CVE-2013-1848
 - "fs/ext3/super.c in the Linux kernel before 3.8.4 uses incorrect arguments to functions in certain circumstances related to printk input, which allows local users to conduct format-string attacks and possibly gain privileges via a crafted application."
 - local users may be able to gain superuser (i.e., root) privileges
- CVE-2016-4071
 - "Format string vulnerability in the php_snmp_error function in ext/snmp/snmp.c in PHP before 5.5.34, 5.6.x before 5.6.20, and 7.x before 7.0.5 allows remote attackers to execute arbitrary code via format string specifiers in an SNMP::get call."
- Attacks may combine vulnerabilities
 - first, compromise a process with restricted privileges (e.g., webserver)
 - then, use privilege escalation to gain root access

Malicious Filename as Input

- Example:
 - CVE-2012-0809: "Format string vulnerability in the sudo_debug function in Sudo 1.8.0 through 1.8.3p1 allows local users to execute arbitrary code via format string sequences in the program name for sudo."
 - for debugging purposes, the name of the sudo program (e.g., /usr/bin/sudo) is used as part of the format string passed to fprintf(...)
- Exploitation
 - user creates symbolic link to sudo:


```
ln -s /usr/bin/sudo
./<<maliciouscode>><<address>>%x%x%n
```
 - user executes sudo (with debugging enabled) using the symbolic link:

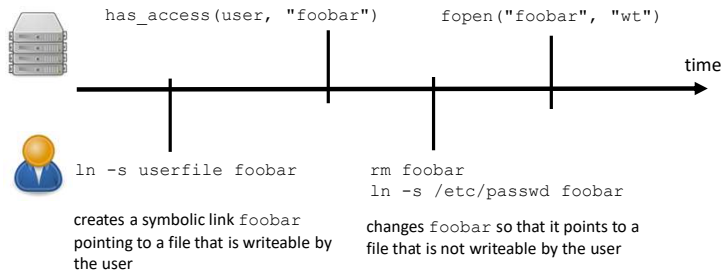

```
./<<malicious code>><<address>>%x%x%n -D9
```
- Note that the exploit does not require the user to be on the sudoers list. The sudoers file is a file Linux and Unix administrators use to allocate system rights to system users.

4. Race Conditions

- Race condition
 - When results depend on the sequence or timing of uncontrollable events.
 - for example, when software output depends on how the OS schedules the execution of multiple processes or threads.
- Typically happens when interacting with
 - memory shared by multiple processes (or threads),
 - file system, or
 - signals and other inter-process communication mechanisms.
- Race condition bugs and errors
 - happen when events do not occur in the intended order
 - typically, very difficult to reproduce and debug

Race Condition Vulnerability

- Process with superuser privileges:
 - takes as input from the user a filename and some data
 - if the user has access to the file, the process writes data to it



Exploiting a Race Condition

- Challenge (for the attacker): time between checking access and opening the file is very short
- Attack approaches
 - try multiple times (if possible)
 - slow down the target process
 - increase computational load on the machine
 - computational complexity attacks

Attack techniques for file paths

- deeply nested directories:


```
filename = "this/is/deeply/nested/.../keep/going/.../wait/for/it/.../almost/there/.../finally/target_file"
```
- chain of symbolic links:


```
ln -s target_file link1; ln -s link1 link2; ...; ln -s linkN filename
```

Preventing Race Conditions

- Time of check to time of use
 - we cannot allow any changes in this interval
 - trying to make it short is not enough
- Prevention techniques
 - work with file descriptors instead of filenames
 - rely on filesystem access checks
 - be careful with directories that are writable by everyone (e.g., /tmp/)
 - lock resources (files, databases, etc.)
 - look out for non-atomic operations (e.g., num++)
 - synchronization (e.g., semaphore, mutex)

Race

- These shared memory accesses may happen concurrently
 - if two different processes execute on separate processors, or
 - asynchronously, when a thread sleeps for some time.
- A common cause of race conditions is when a program uses static variables.
- When accessing shared memory locations without appropriate synchronization, ensuring that only one line can utilize them at a time is critical.
- A race condition can occur while accessing a file: the adversary can trick the system by replacing a file with their version and cause the system to read the malicious file.

Race Condition Vulnerability

- CVE-2014-0196
 - “The `n_tty_write` function in `drivers/tty/n_tty.c` in the Linux kernel through 3.14.3 does not properly manage tty driver access in the “LECHO & !OPOST” case, which allows local users to cause a denial of service (memory corruption and system crash) or gain privileges by triggering a race condition involving read and write operations with long strings.”

5. Buffer Overflow Countermeasures

Compile-time hardening new software	Run-time protecting existing software
--	--

- | | |
|--|---|
| <ul style="list-style-type: none"> • programming languages • safe functions and libraries • compiler extensions | <ul style="list-style-type: none"> • executable space protection • address space layout randomization |
|--|---|

unpredictability

5.1 Compile-Time Prevention

Programming Languages

- Programming languages and platforms that do not allow direct memory access typically prevent buffer overflows
 - C# and other managed .NET languages
 - Java, Python, PHP, Perl
- However,
 - safety and security may come at the cost of lower performance (bounds checking can increase execution time)
 - native code written in other languages is still vulnerable (e.g., through Java Native Interface)
 - be careful with libraries that might rely on native code
 - there might be vulnerabilities in the interpreter/virtual machine of the language (e.g., JVM or Common Language Runtime CLR)

Avoid UnSafe Functions

- Unsafe
 - `strcpy(char *src, char *dst)`
 - `strcat(char *s1, char *s2)`
 - `sprintf(char *str, [1]sep char* format, ...)`
 - `gets(char *str)`
- Less Unsafe
 - `strncpy(char *src, char *dst, size_t n)`
 - `strncpy(char *src, char *dst, size_t n)`
 - `strncat(char *s1, char *s2, size_t n)`
 - `strlcat(char *s1, char *s2, size_t n)`
 - `snprintf(char *str, size_t n, char* format, ...)`
 - `fgets(char *str, int n, FILE *file)`
(file can be stdin)

Use Safe Functions

- Safe user input:

```
char buffer[SIZE];
fgets(buffer, sizeof(buffer), stdin);
```

- Memory copying (`memcpy` and `memmove`)

- less error-prone since we have to specify the source size
- `memcpy_s(void *dst, size_t dn, void *src, size_t sn)`

Safe Libraries

- C++ Standard Library

- String class `std::string`

- mostly safe:

```
str3 = str1 + str2
```

- however, we can shoot ourselves in the foot with the `[]` operator:

```
string str("ouch");
str[100] = '!';
```

- Array class `std::vector`

- dynamic-size array implementation
- but we can use it in an unsafe manner:

```
std::vector<int> array(4);
array[10] = 1;
```

Stack Canary Compiler Extension



In the coal mines, when the bird stopped singing - the miners evacuated immediately.



We need a canary to warn us when the return address has been overwritten.

local variables of caller

function parameters

return address

saved frame pointer

local variables of function
(e.g., `char buffer[16]`)

Stack Canary Compiler Extension

- At the beginning of the function, a special value is placed on the stack between the local variables and the return address
- Buffer overflows that reach the return address will necessarily overwrite the special value
- Before returning, the function checks if the value is intact



- Stack-based buffer overflow attacks can be detected

buffer overflow

local variables of caller

function parameters

return address

canary

saved frame pointer

local variables of function
(e.g., `char buffer[16]`)

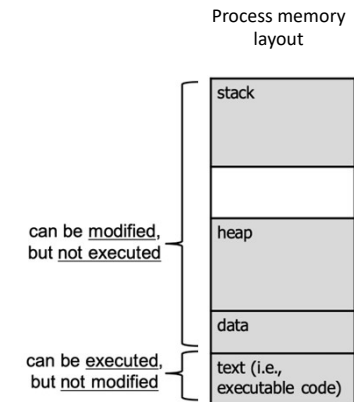
Canary Values & Implementations

- Typical canary values
 - Terminator: contains zero bytes, newline (CR, LF), etc.
 - if input cannot contain one of these values (*e.g.*, strings cannot contain zero bytes), then the canary value is necessarily modified by the exploit
 - Random: random value, typically chosen when the program starts
 - random XOR: canary is the XOR of the random value and the return address
- Implementations
 - Microsoft Visual Studio: `/GS` (Buffer Security Check) option
 - enabled by default, protects control data by creating a copy
 - GCC (GNU Compiler Collection): `-fstack-protector` flag
 - enabled by default, based on a random value

5.2 Run-Time Countermeasures

Executable Space Protection

- Lot of exploits build on injecting and executing malicious code.
- By separating the memory space of a process into executable and modifiable parts, code injection can be prevented.
- *Problem:*
Modern computer architectures do not separate code from data.



Hardware-Based Solution NX Bit

- NX (No-eXecute) bit
 - technology in CPUs for separating code from data
 - each page table entry (i.e., data used for managing a part of the memory by the virtual memory system) has an NX bit
 - 0: code can be executed by CPU
 - 1: code cannot be executed by the CPU
 - Hardware-enforced but needs OS support
- Implementations
 - x86: AMD Enhanced Virus Protection or Intel XD (eXecute Disable) bit (Windows and Linux support it)
 - ARM: XN (eXecute Never) bit
- *Limitation:* cannot fully protect programs that create and execute code at runtime (*e.g.*, just-in-time compilers)
 - for example, web browsers that support JavaScript are at risk

Circumventing Executable Space Protection

- An attacker can re-use existing code from the memory space of the process for malicious purposes.
- Return-to-libc attack
 - for most processes, the standard C library is loaded into memory
 - attacker can change the return address of a function to point to the beginning of a function in the C library
 - common target: system function
 - takes as argument a string, and executes it as a system command with the privileges of the process
 - attacker has control over the stack
 - attacker can set up parameters for the C library function

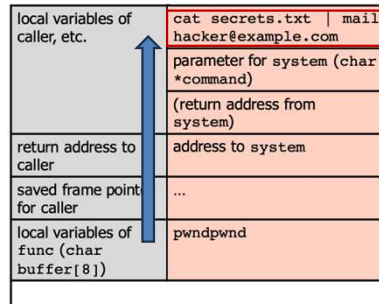
Return-to-libc Attack

Attack string:

pwndpwnd<...><address to system><...><address to command>cat secrets...

```
int system(char *command) {
    // standard C library
    // executes command
    ...
}

void func() {
    char buffer[8];
    gets(buffer);
}
```




Countermeasure

Address Space Layout Randomization

- In order to reliably jump to an exploited code, the attacker needs to know its address
- Address Space Layout Randomization (ASLR)
 - randomly arrange the positions of the executable, the stack, and the heap in the process's address space
 - may prevent return-to-libc attacks
 - most operating systems (e.g., Windows, Linux) implement some randomization
- Counter-countermeasures
 - information leakage (e.g., printf vulnerability)
 - random guessing

ASLR Process Memory Layout

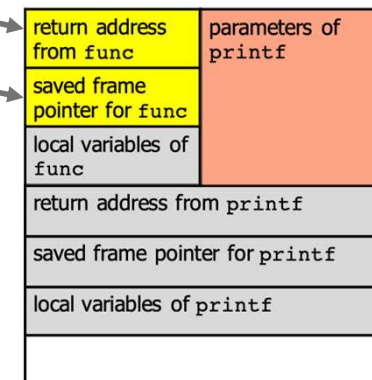


CCM: Information Leakage with printf

- can be used to figure out the address of the executable
- can be used to figure out the address of the stack

```
void func() {
    char test[32];
    gets(test);
    printf("You wrote: ");
    printf(test);
    printf("\n");
}
```

```
$ ./program
%x %x...
You wrote: a4e320ff f0 ...
```



CCM: Random Guessing

- Limitations of ASLR
 - stack and heap cannot be located at any address
 - *example:*
stack might need to be aligned to 16 bytes and
heap might need to be aligned to 4096 bytes
 - on a 32-bit system, a brute-force may be viable
- Heap spraying
 - fill up the memory with a certain sequence of bytes
 - *example:* malicious website trying to compromise the client's web browser might fill up the memory using Javascript
 - *example sequences:*
 - shellcode preceded by "NOP slide"
 - "/////...//bin/sh" for a return-to-libc attack

Next

- Other Vulnerabilities
- Web Vulnerabilities
- Malware